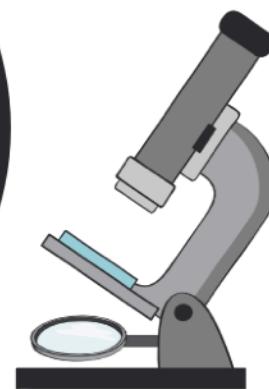




DYNAMIC INSTRUMENTATION FOR GO



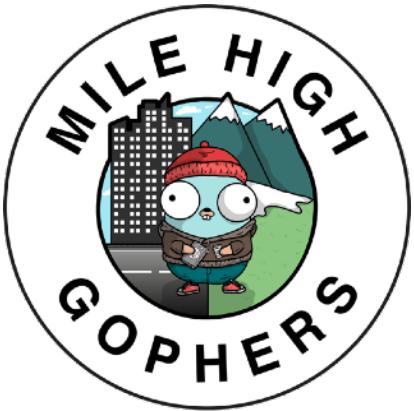


Hello!

I'm Jason Keene

Colorado Gophers
Pivotal
Coinbase

Pivotal
coinbase



Hello!

I'm Jason Keene

Colorado Gophers
Pivotal
Coinbase

BPF

uprobes

ptrace

delve

Instrumentation



Instrumentation is anything that allows you to make measurements





**It's
Story
Time**

Antonie van Leeuwenhoek



Dutch cloth merchant in the 1600s

He wanted a way to inspect the fibers of clothing he sold

He improved the design of the early microscope using lenses that allowed him to see clear, focused images of high magnification



He used his microscope to study the water in Holland.

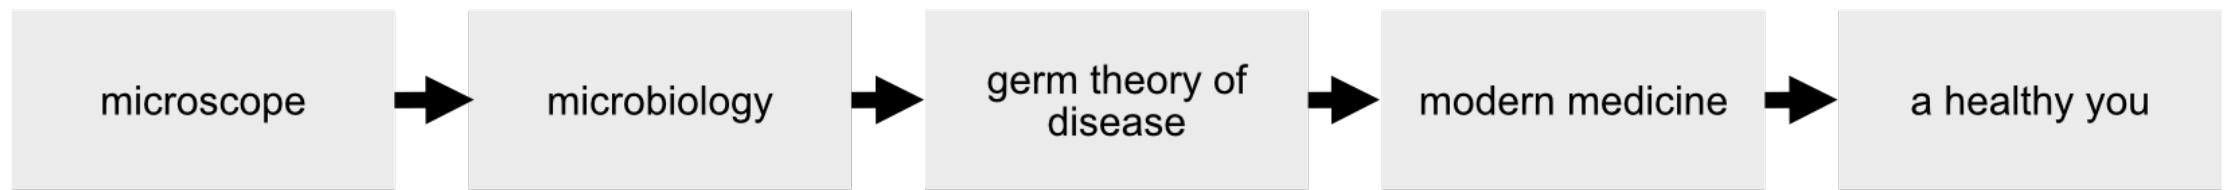
He discovered that the water contained little creatures,
single-celled organisms that he called

Animalcules (Bacteria)



So why am I telling you about this?

The development of the microscope was required to make progress in biology.

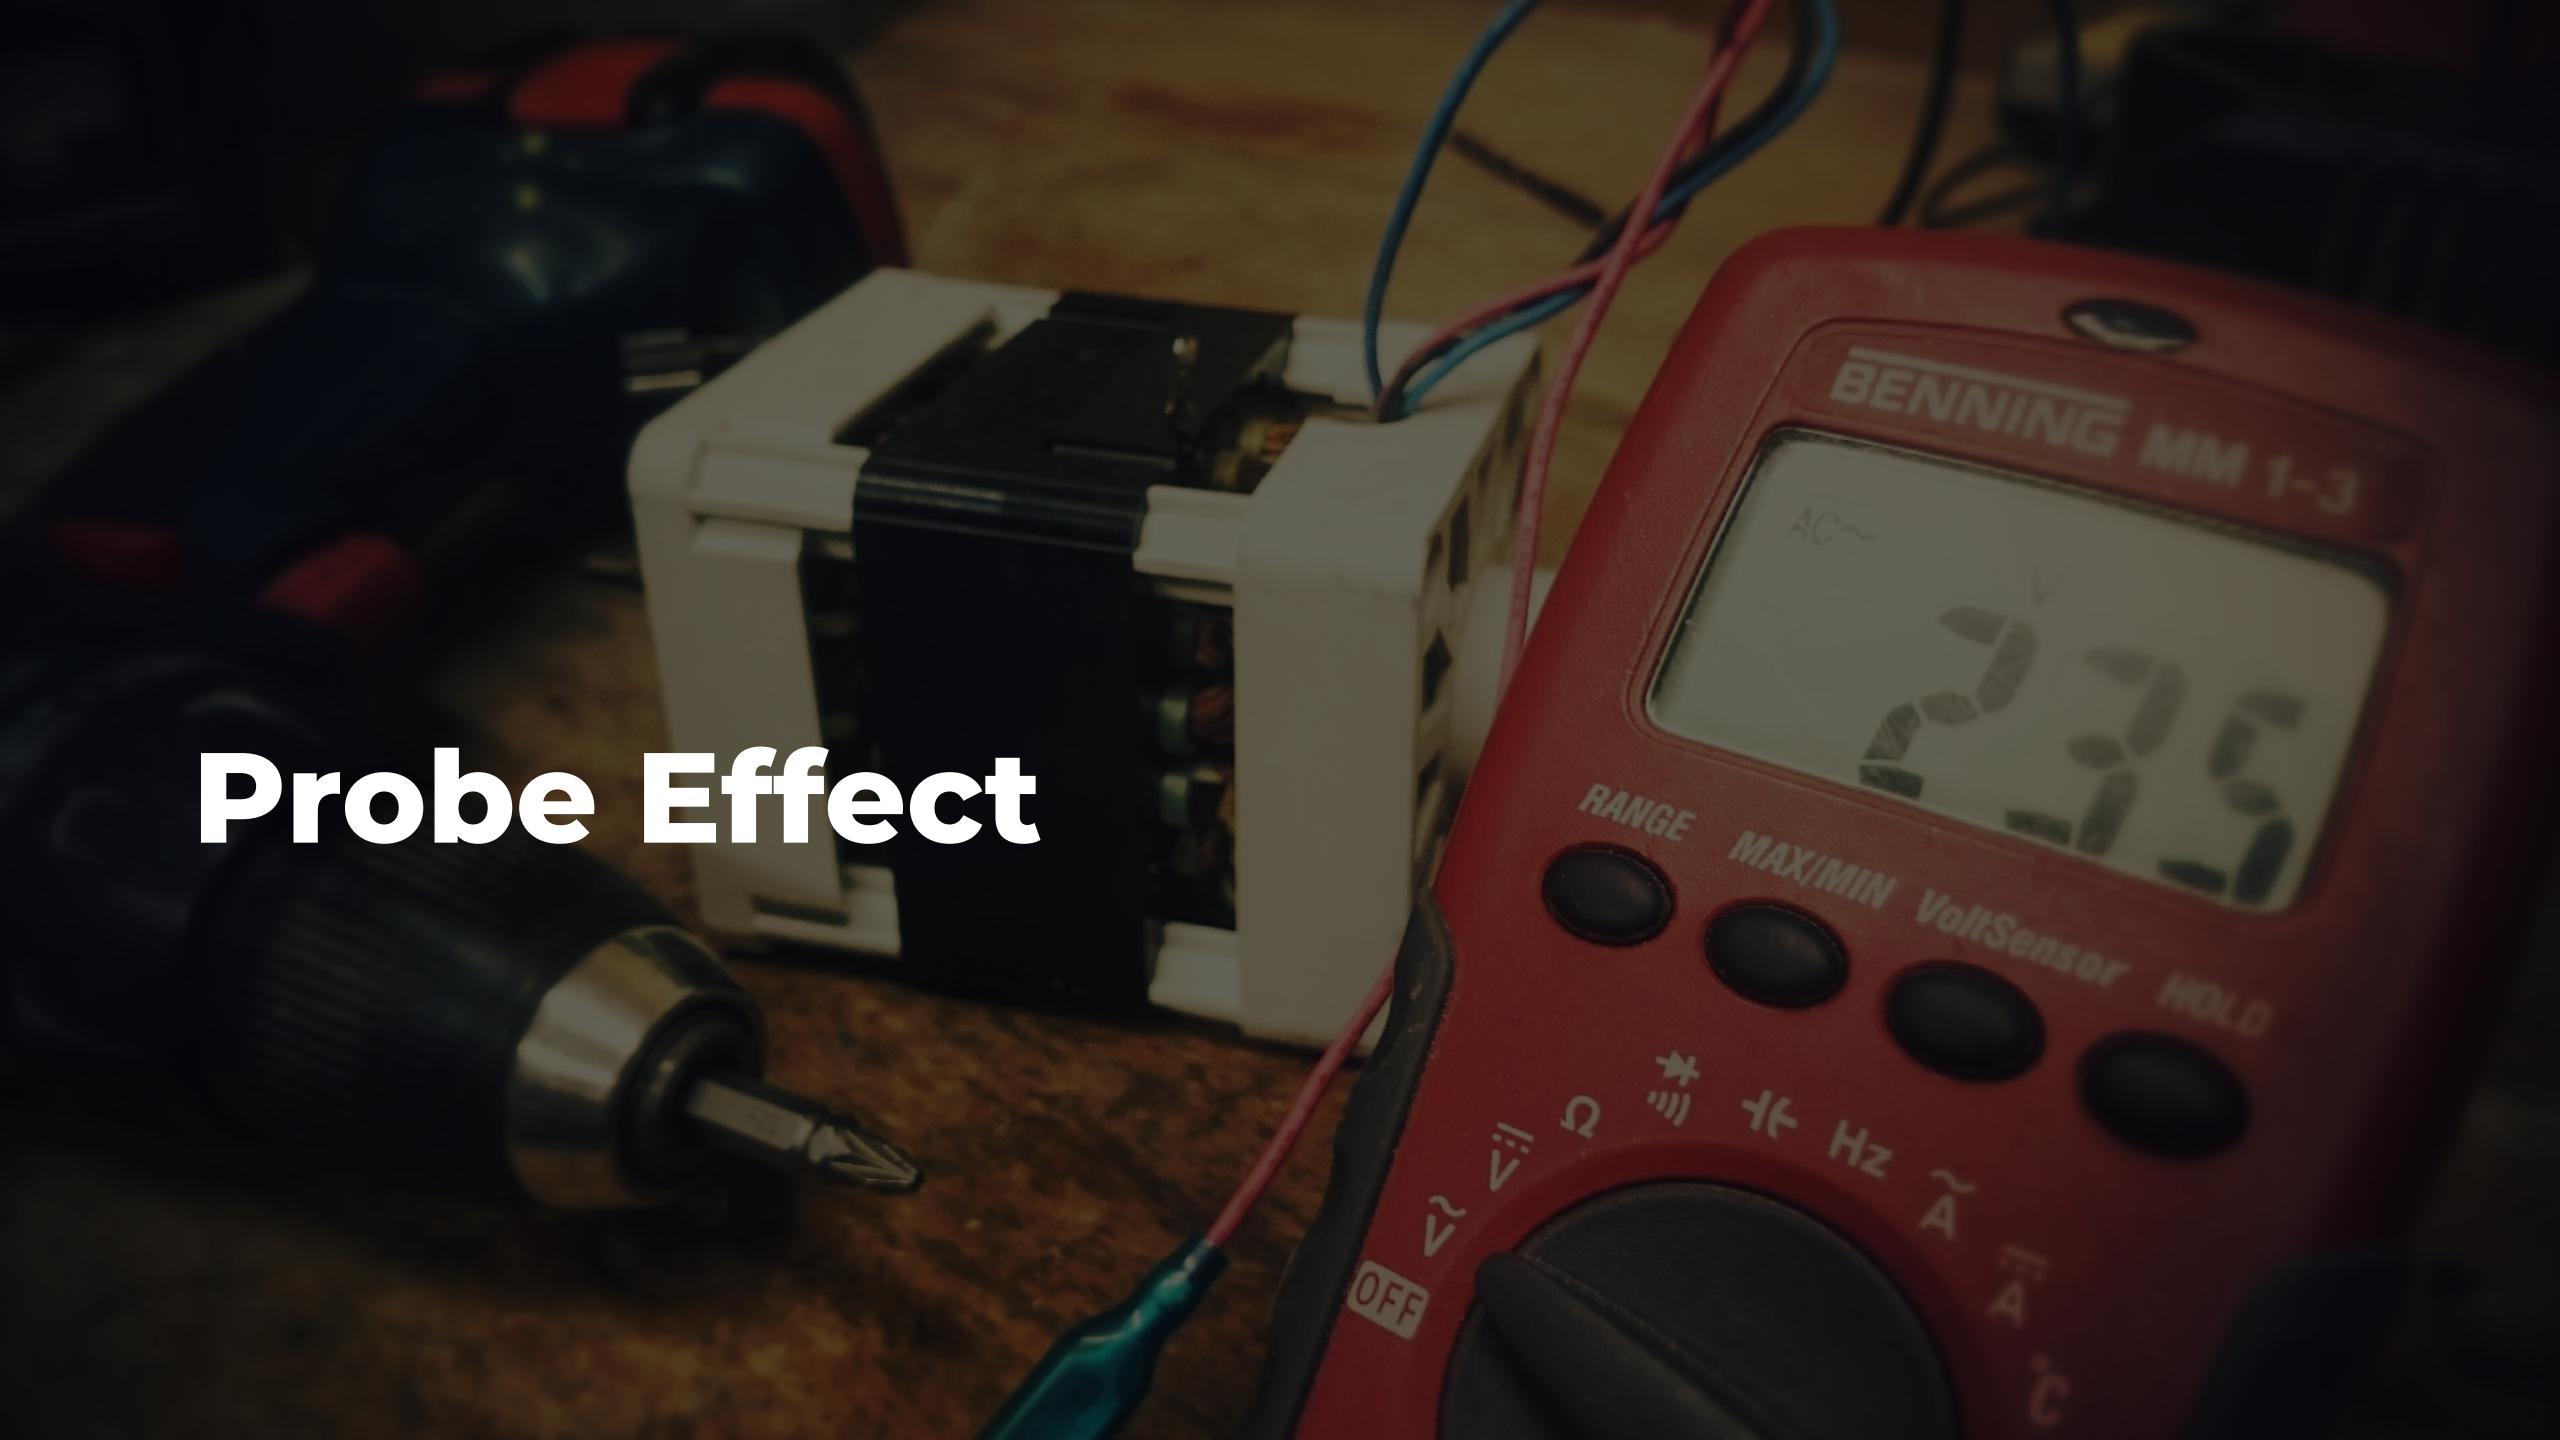


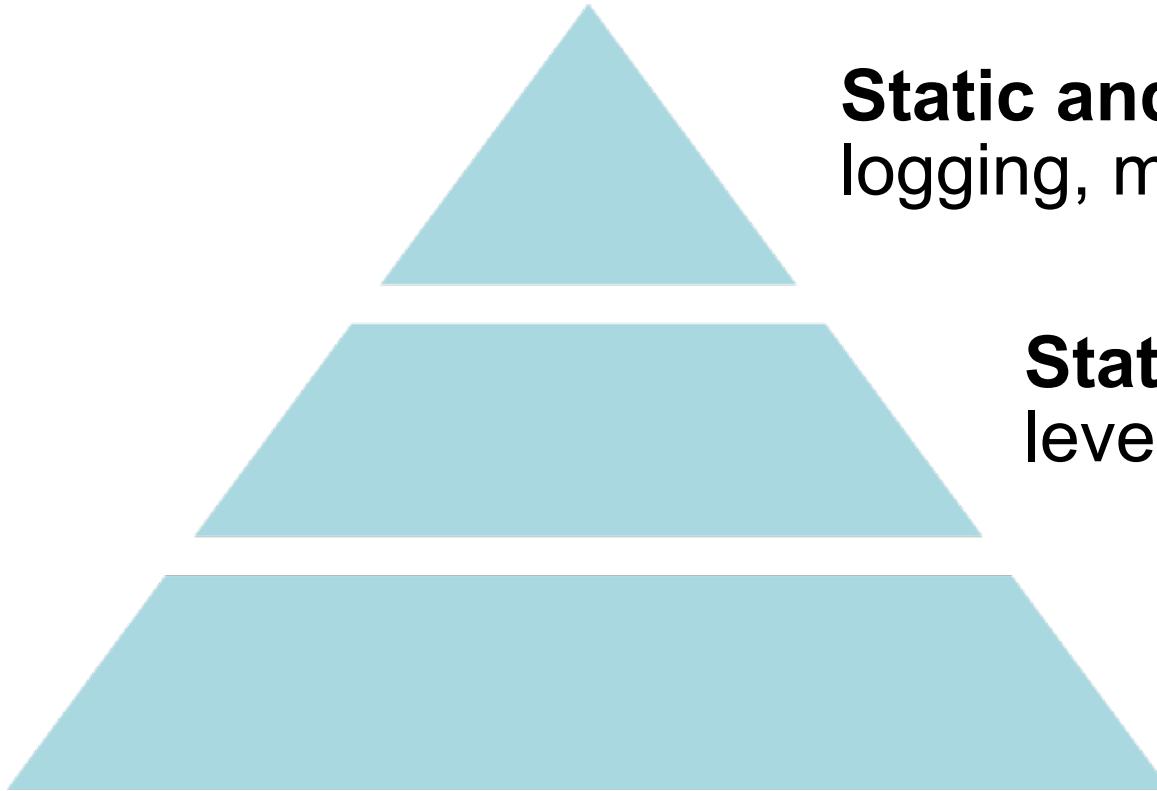
A new kind of instrumentation was required to achieve a new level of understanding in biology.

Instrumentation in Software



Probe Effect



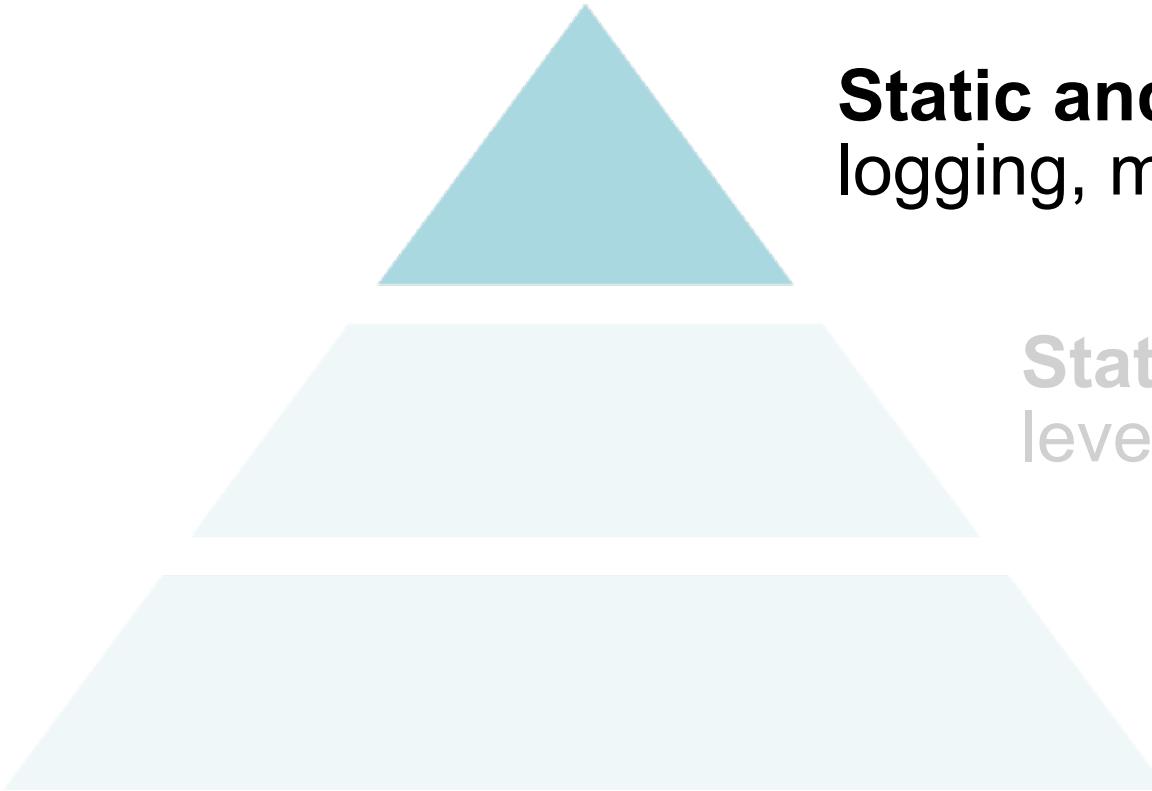


Measurement Coverage

Static and Always On
logging, metrics, distributed tracing

Static and Requires Activation
leveled logging, pprof, usdt probes

Dynamic Instrumentation
debuggers, tracers, uprobes/bpf

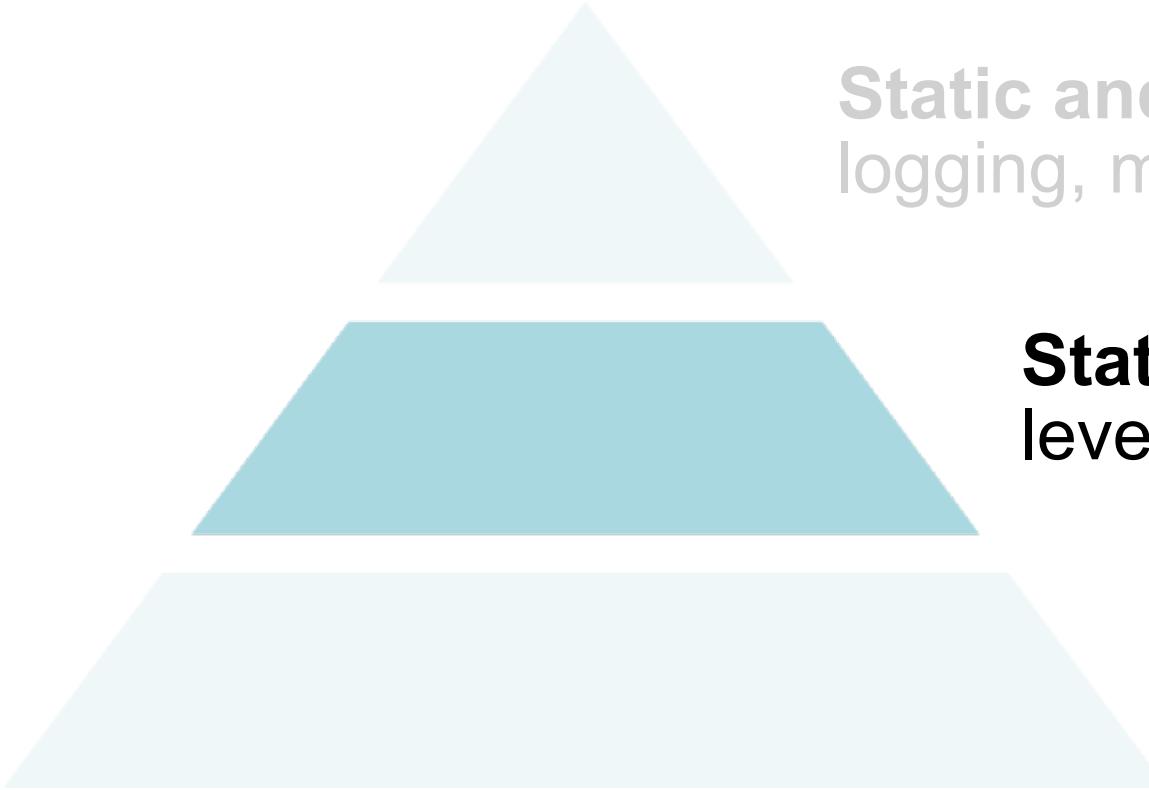


Static and Always On
logging, metrics, distributed tracing

Static and Requires Activation
leveled logging, pprof, usdt probes

Dynamic Instrumentation
debuggers, tracers, uprobes/bpf

Measurement Coverage



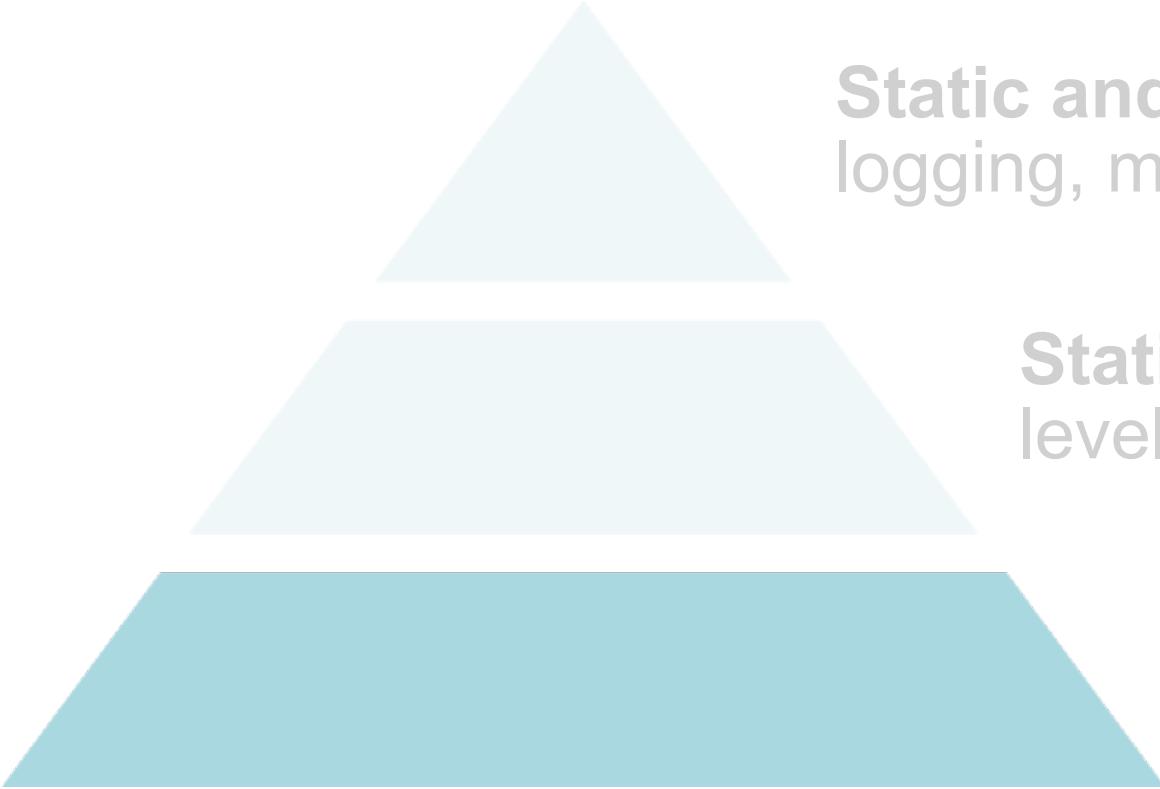
Static and Always On
logging, metrics, distributed tracing

Static and Requires Activation
leveled logging, pprof, usdt probes

Dynamic Instrumentation
debuggers, tracers, uprobes/bpf

Measurement Coverage

Measurement Coverage



Static and Always On
logging, metrics, distributed tracing

Static and Requires Activation
leveled logging, pprof, usdt probes

Dynamic Instrumentation
debuggers, tracers, uprobes/bpf

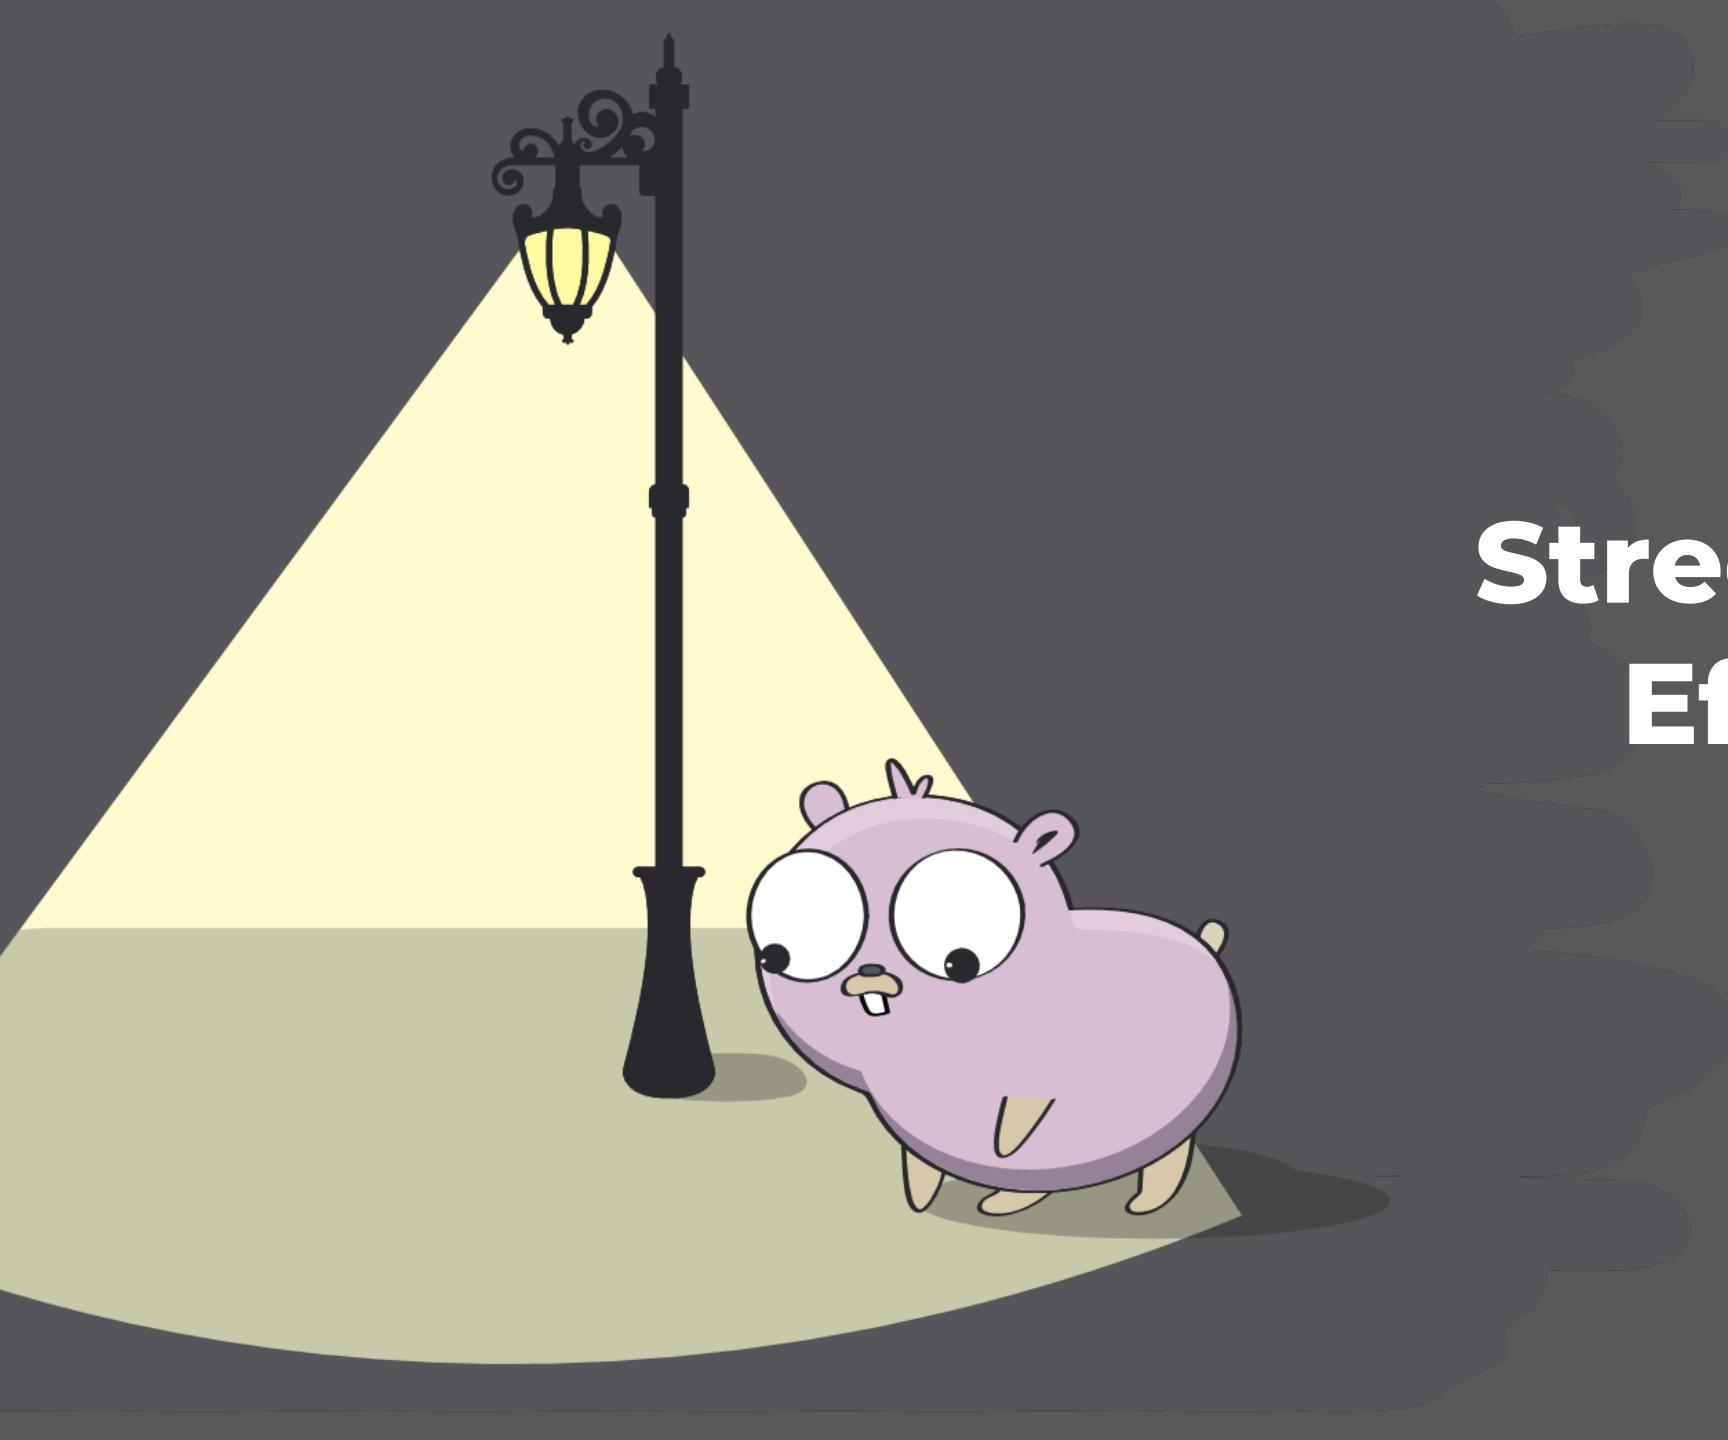
We want tools that can answer arbitrary questions about our software

Intercept any point of execution

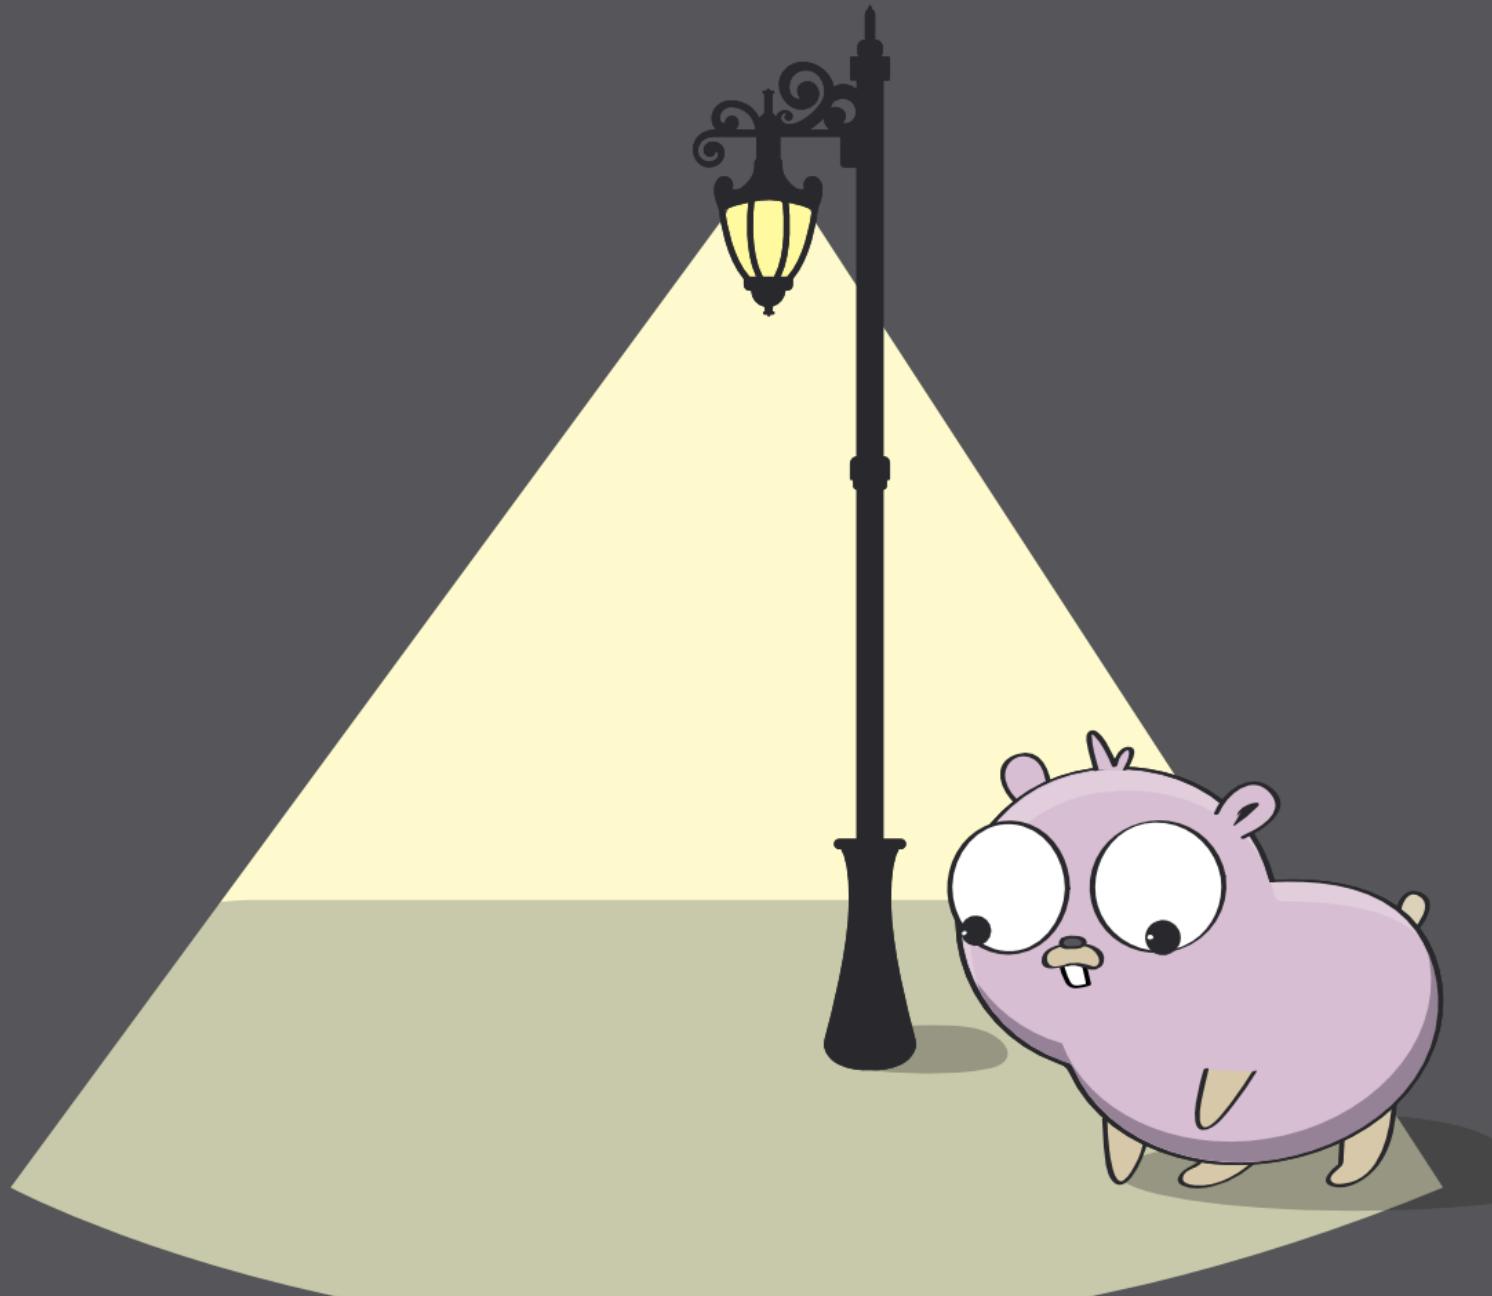
Without restarting the process

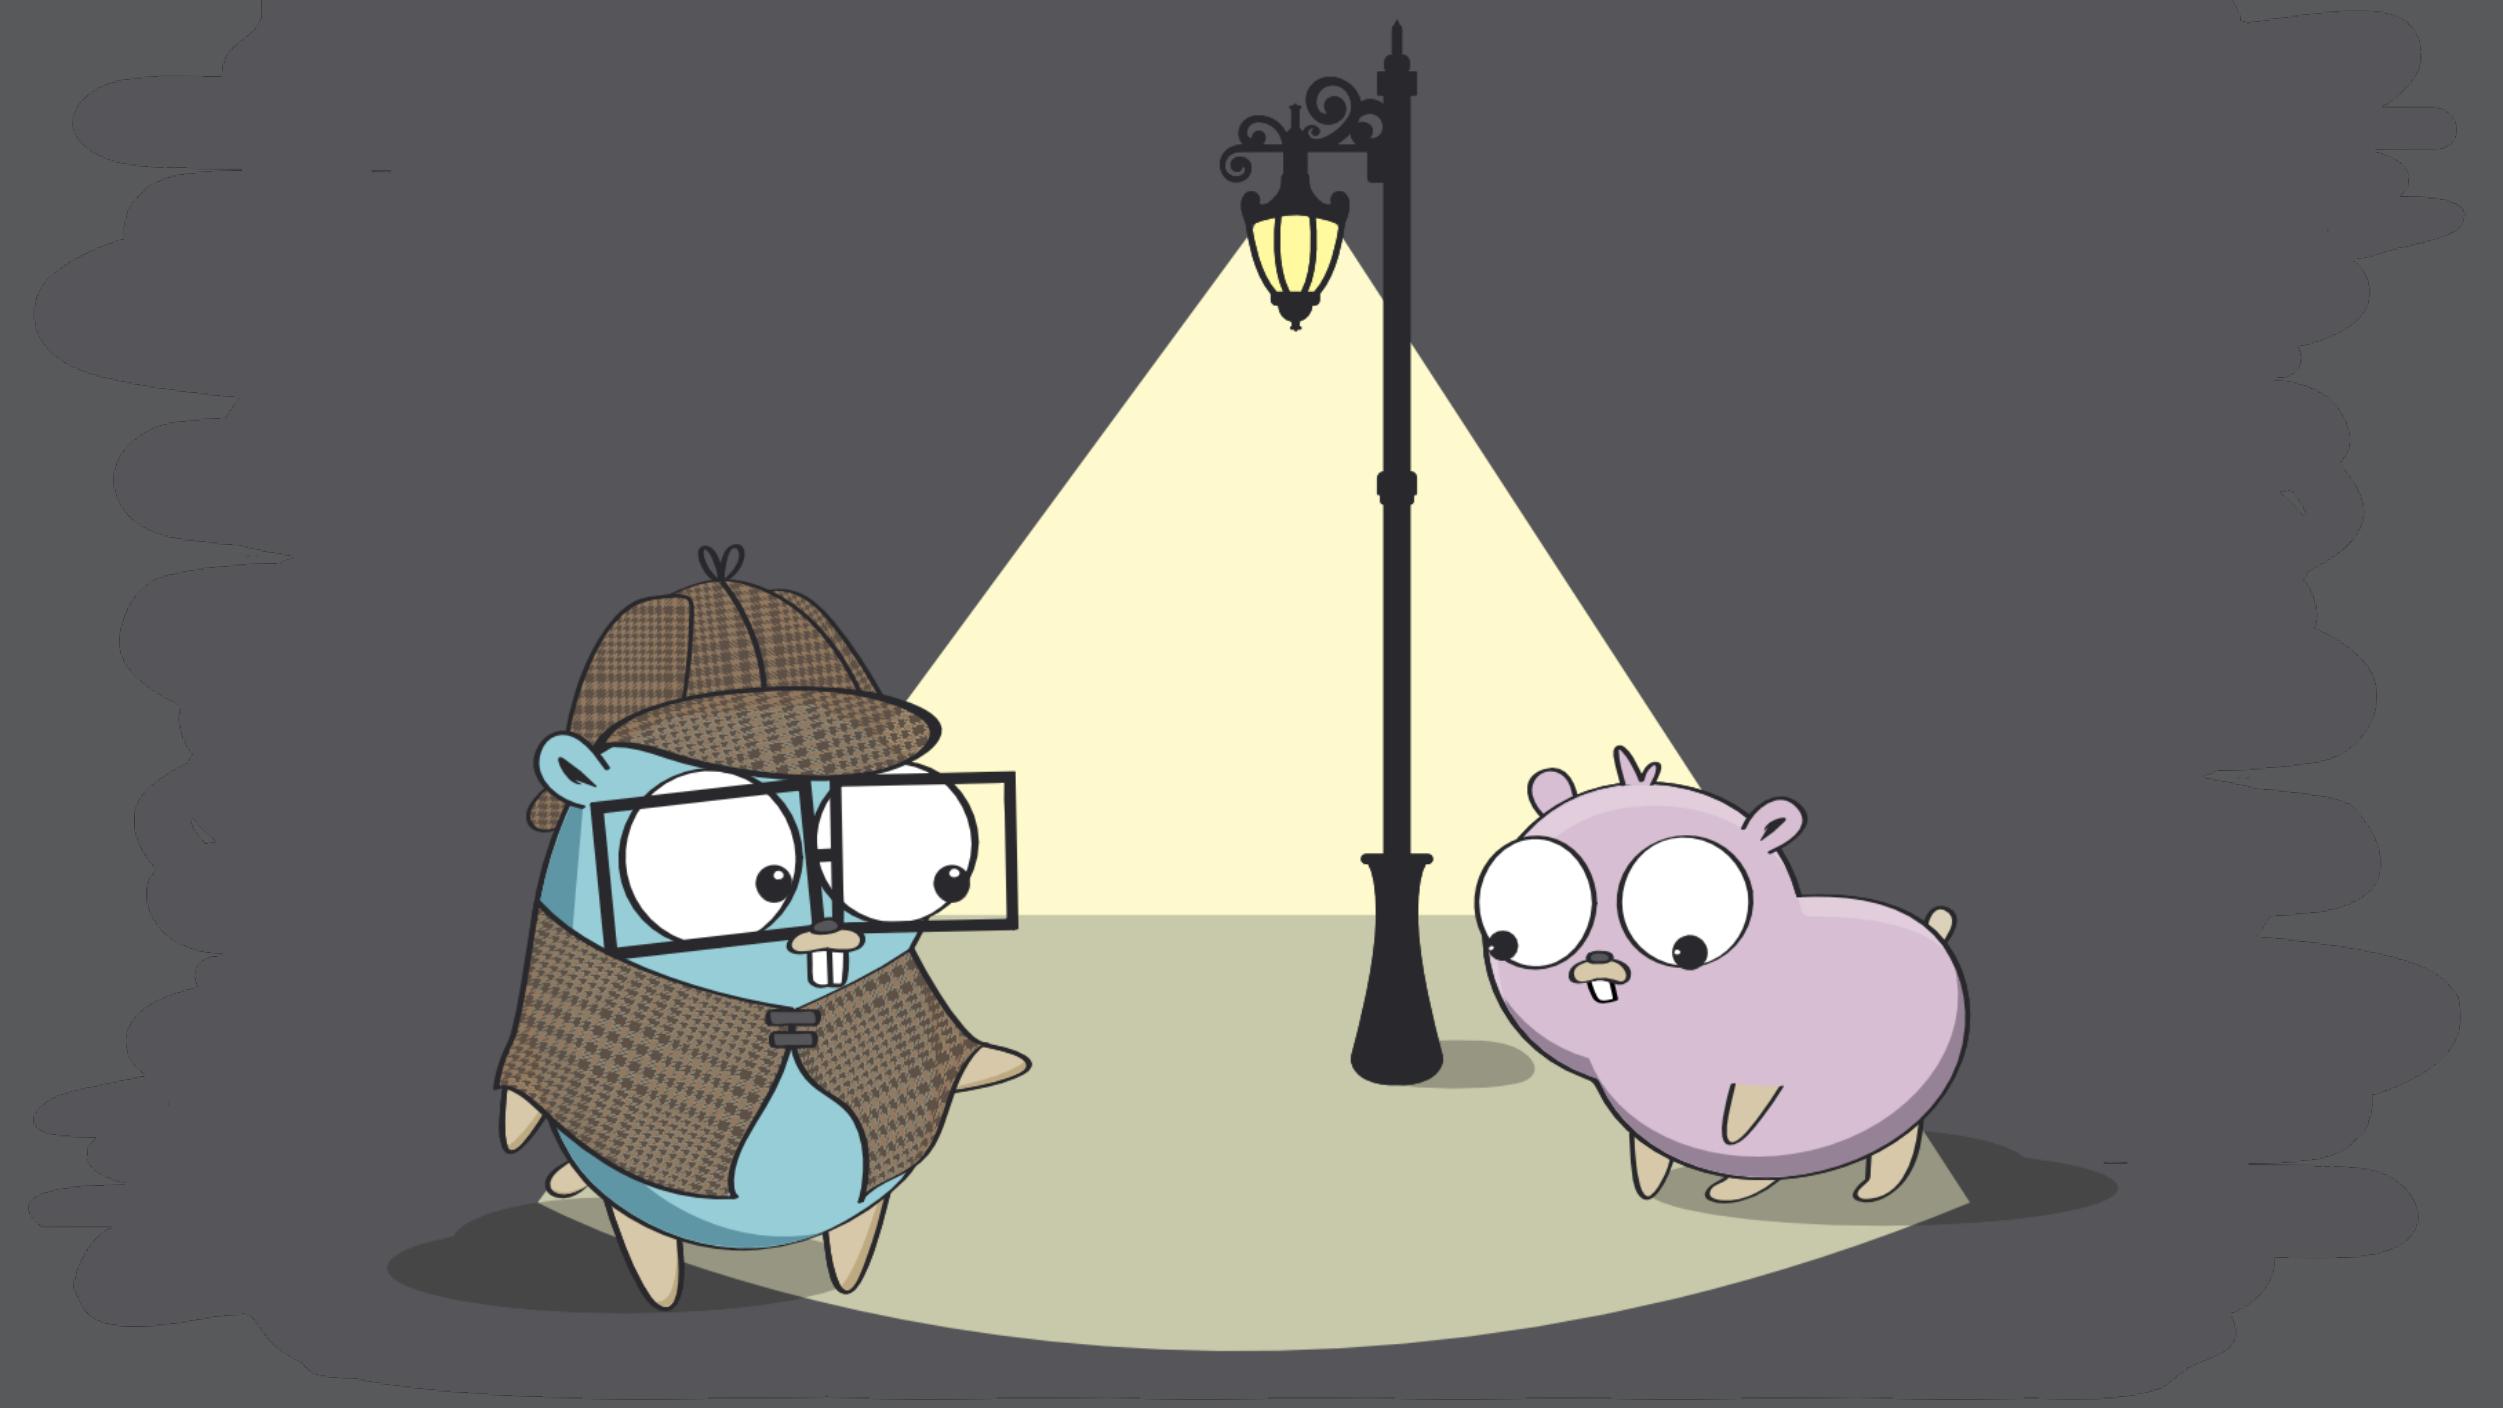
With low overhead

And do it all safely

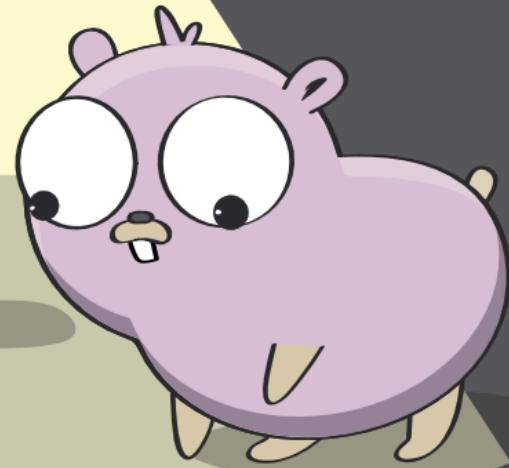
A cartoon illustration of a pink gopher character with large white eyes and a small brown nose. It is standing on a light-colored ground surface, looking up at a black street lamp post. The lamp's light beam is a bright yellow triangle pointing towards the gopher. The background is dark gray with some abstract cloud shapes.

Streetlight Effect





this is where you
lost your wallet?



this is where you
lost your wallet?

no, I lost it in the park.
but this is where
the light is.

Dynamic Instrumentation
allows you to look in the park.

I want a toolkit that can examine running programs and answer questions about them. Questions I don't know until the problem arises...

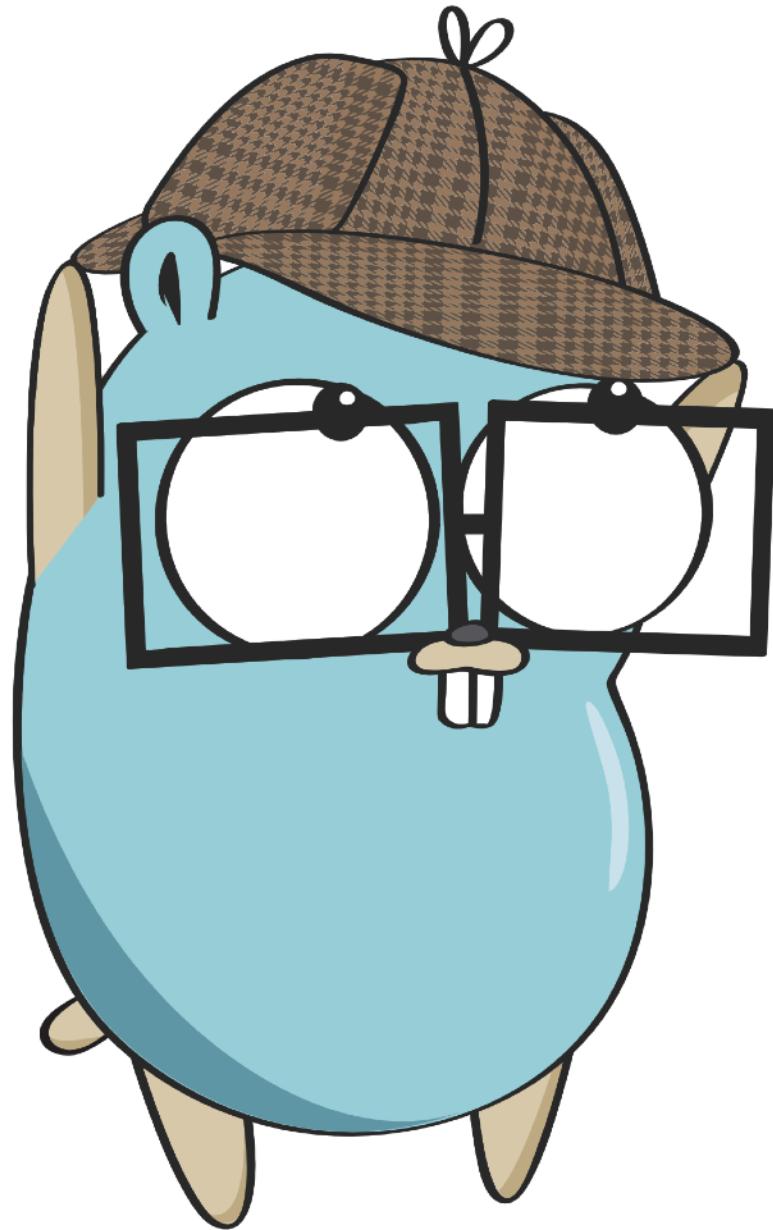
I'd like not to require compilation-time changes. I'd like to come to an arbitrary program while it's running...

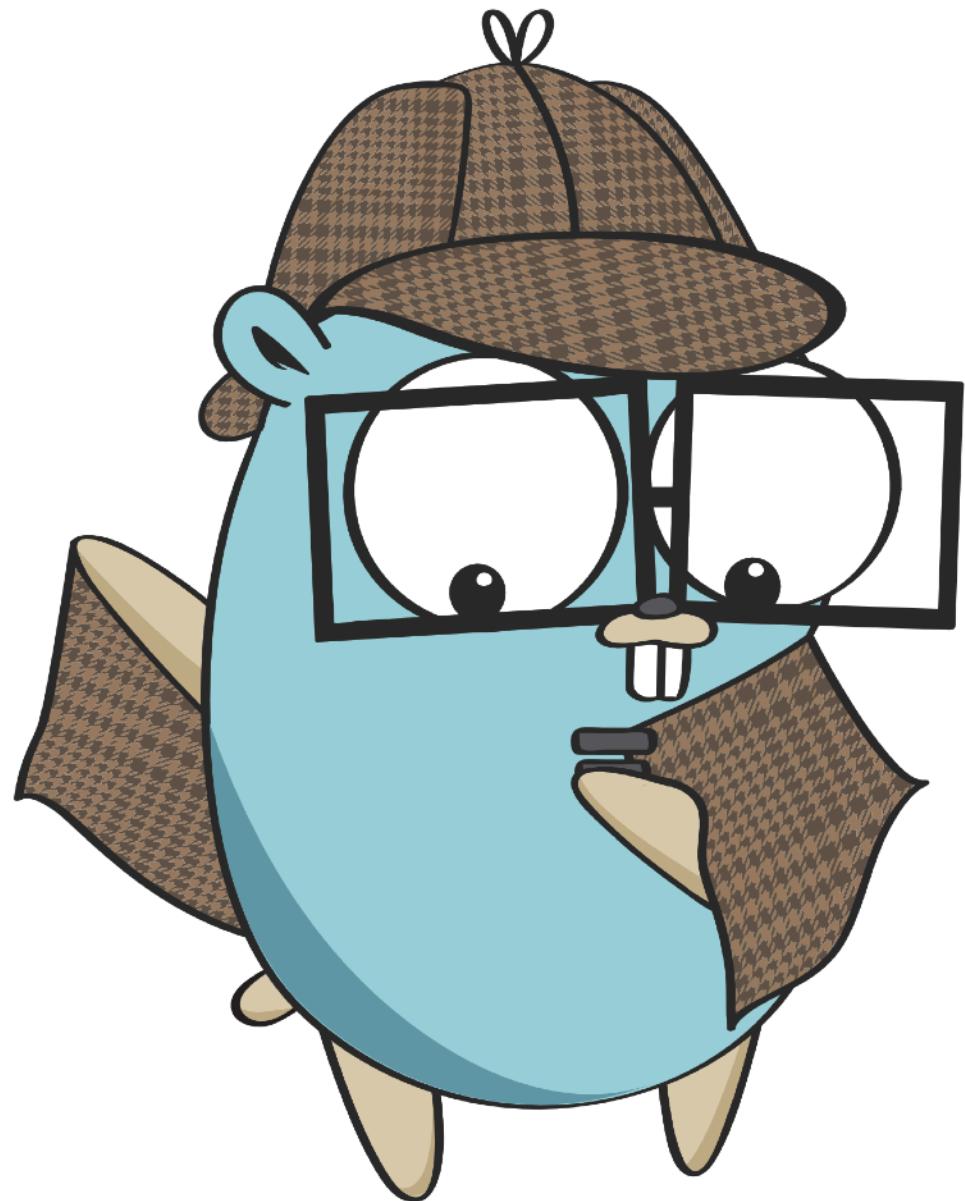
I want a programming interface above which I can build whatever is needed, either generally-useful tools or ad hoc probes. And the programming language I want to use for this is Go.

— Rob Pike

What Rob wants does not exist, yet...

However, there have been many advances
that provide a good deal of capability.









ptrace and Debuggers

Debuggers are Incredibly
Powerful Tools

A traditional Japanese woodblock-style illustration of a landscape. In the foreground, there are stylized waves and clouds. Several small boats with figures are scattered across the water. A large, dark, textured shape, possibly a rock or a stylized tree trunk, is visible on the left. The background shows more distant waves and a light-colored sky.

Starting GDB

Debuggers Use the
ptrace System Call

- Suspend/Continue Execution
- Single-step
- Read and Write to Memory
- Read and Write Registers
- Intercept Signals

Let's see how debuggers use ptrace to set breakpoints and intercept execution

USER

KERNEL

Go Process

Interrupt
Handler

Delve

ptrace

**Let's Take a Look at
the Instructions as the
Program Runs**

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace

B8
00

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace

**Delve Uses ptrace to
Set Breakpoints**

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

**When a Breakpoint is
Hit Delve Takes Over**

00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

CC

00

B8

USER

KERNEL

Go Process

Delve

Interrupt
Handler

ptrace



00

B8

01

00

CC

00

B8

USER

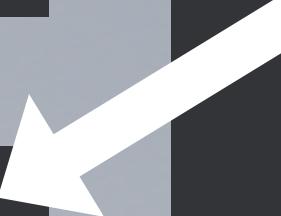
KERNEL

Go Process

Delve

Interrupt
Handler

ptrace



00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve



KERNEL

Interrupt
Handler

ptrace

**Eventually Delve uses
ptrace to Resume Execution**

00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

B8
Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace



00

B8

01

00

CC

00

B8

USER

Go Process

B8
Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve

KERNEL

Interrupt
Handler

ptrace

The major problem with this is suspended execution.

00

B8

01

00

CC

00

B8

USER

Go Process

B8

Delve



KERNEL

Interrupt
Handler

ptrace

We can minimize the amount of time the debugger is running by automating our interactions with it.

Delve exposes a JSON-RPC API that we can use!

We can write a program that uses delve's API to add the instrumentation we want in an automated way!

00

B8

01

00

B8

00

B8

USER

Go Process

Delve

KERNEL

Interrupt
Handler

ptrace

00

B8

01

00

B8

00

B8

USER

Go Process

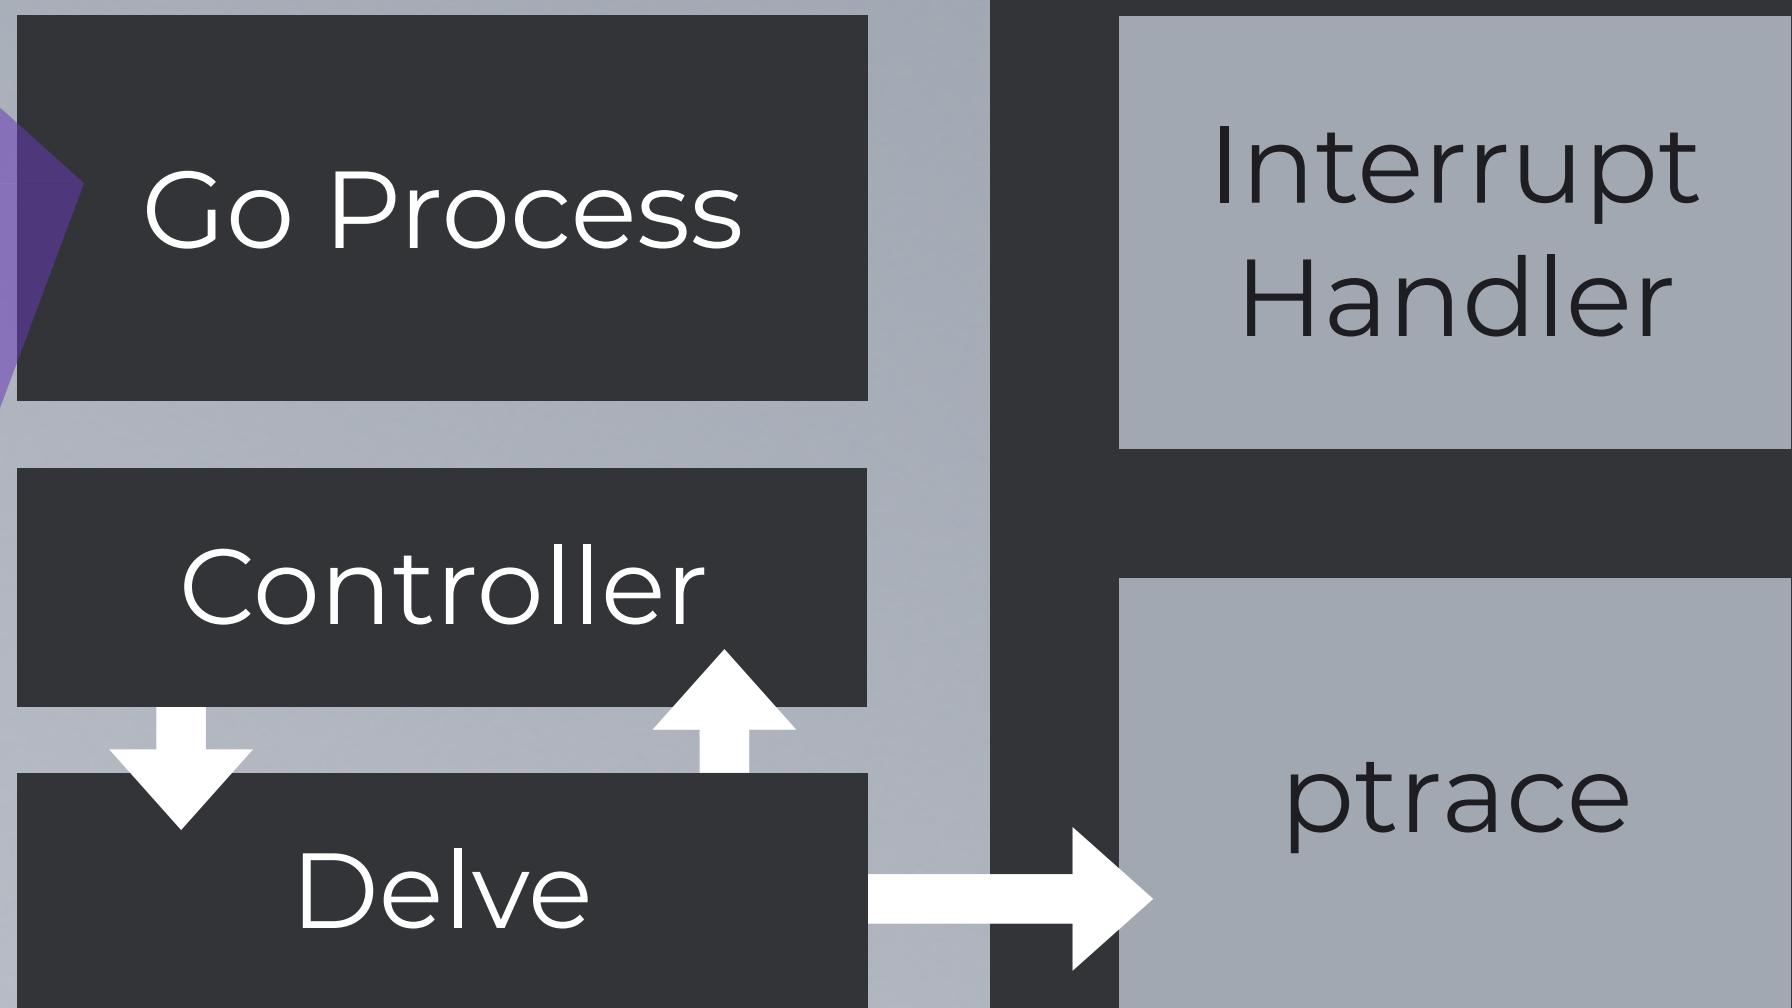
Controller

Delve

KERNEL

Interrupt
Handler

ptrace



Let's first create a test program that requires some dynamic instrumentation.

```
var counter int64

func doWork() {
    counter++
}

func main() {
    fmt.Println("This is a blackbox. Read my counter.")
    for {
        doWork()
    }
}
```

```
$ /tmp/counter
```

```
dlv := exec.Command(  
    "dlv",  
    "attach",  
    strconv.Itoa(pid),  
    "--headless",  
    "--accept-multiprocessor",  
    "--api-version=2",  
    "--listen="+addr,  
)  
dlv.Start()  
waitForUp(addr)
```

```
func waitToBeUp(addr string) {
    done := time.After(5 * time.Second)
    for {
        conn, err := net.DialTimeout("tcp", addr, 2*time.Second)
        if err == nil {
            conn.Close()
            return
        }
        select {
        case <-done:
            log.Fatalf("Stuck waiting for %q to be up: %s", addr, err)
        default:
        }
    }
}
```

```
client := rpc2.NewClient(addr)
client.Continue()
defer func() {
    client.Halt()
    client.Detach(false)
}()
```

```
locations, _ := client.FindLocation(api.EvalScope{  
    GoroutineID: -1,  
, "main.doWork")  
pc := locations[0].PC
```

```
bp := &api.Breakpoint{  
    Name:      "myAwesomeBreakpoint",  
    Addr:      pc,  
    Variables: []string{"counter"},  
}
```

```
for {
    client.Halt()
    client.CreateBreakpoint(bp)
    state := <-client.Continue()
    client.ClearBreakpointByName(bp.Name)
    client.Continue()

    counter := state.currentThread.BreakpointInfo.Variables[0].Value
    fmt.Printf("counter: %s", counter)

    time.Sleep(time.Second)
}
```

```
$ ./tmp/counter --report
```

{

```
$ ./tmp/lowhz --pid $(pgrep counter)
```

What happens if we intercept
every call to doWork?

```
var sharedCounter int64
go func() {
    for {
        fmt.Printf("counter: %d", atomic.LoadInt64(&sharedCounter))
        time.Sleep(time.Second)
    }
}()

client.CreateBreakpoint(bp)
stateCh := client.Continue()

for {
    state := <-stateCh
    stateCh = client.Continue()

    counterStr := state.currentThread.BreakpointInfo.Variables[0].Value
    counter, _ := strconv.ParseInt(counterStr, 10, 64)
    atomic.StoreInt64(&sharedCounter, counter)
}
```

```
$ /tmp/counter --report
```

} {

```
$ ./tmp/highhz --pid $(pgrep counter)
```

Why are High Frequency
Breakpoints So Bad?

00

USER

KERNEL

B8

Lots of IO for State Transfer

upt
er

01

00

CC

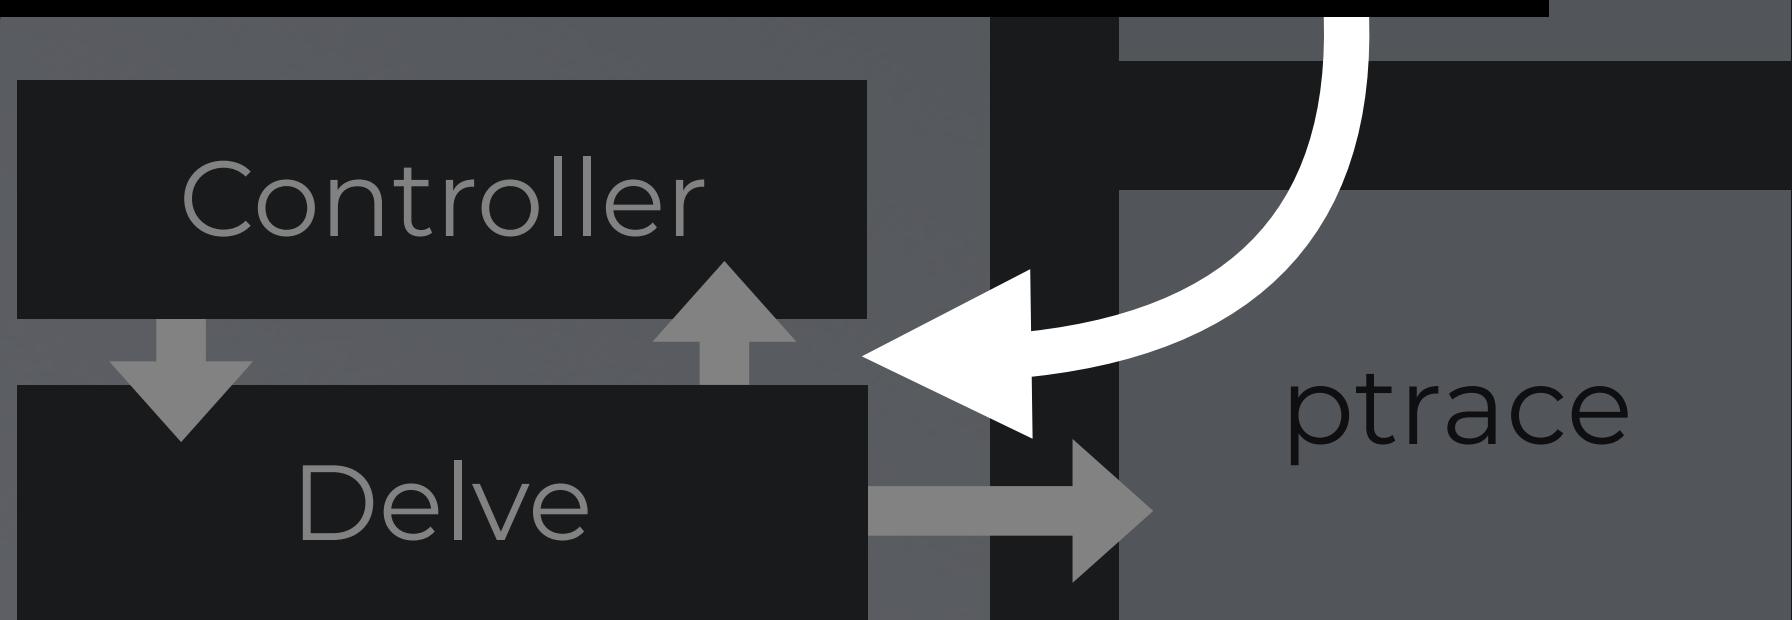
00

B8

Controller

Delve

ptrace



00

USER

KERNEL

Lots of Context
Switches

00

CC

00

B8

:ess

B8

Delve

Interrupt
Handler

ptrace



00

B8

USER

KERNEL



Interrupt

Acessing the process's memory and registers using ptrace can be 10x to 100x slower and incurs several context switches

00

B8

B8

Delve

ptrace

Can We Do Better?

00

B8

01

00

B8

00

B8

USER

Go Process

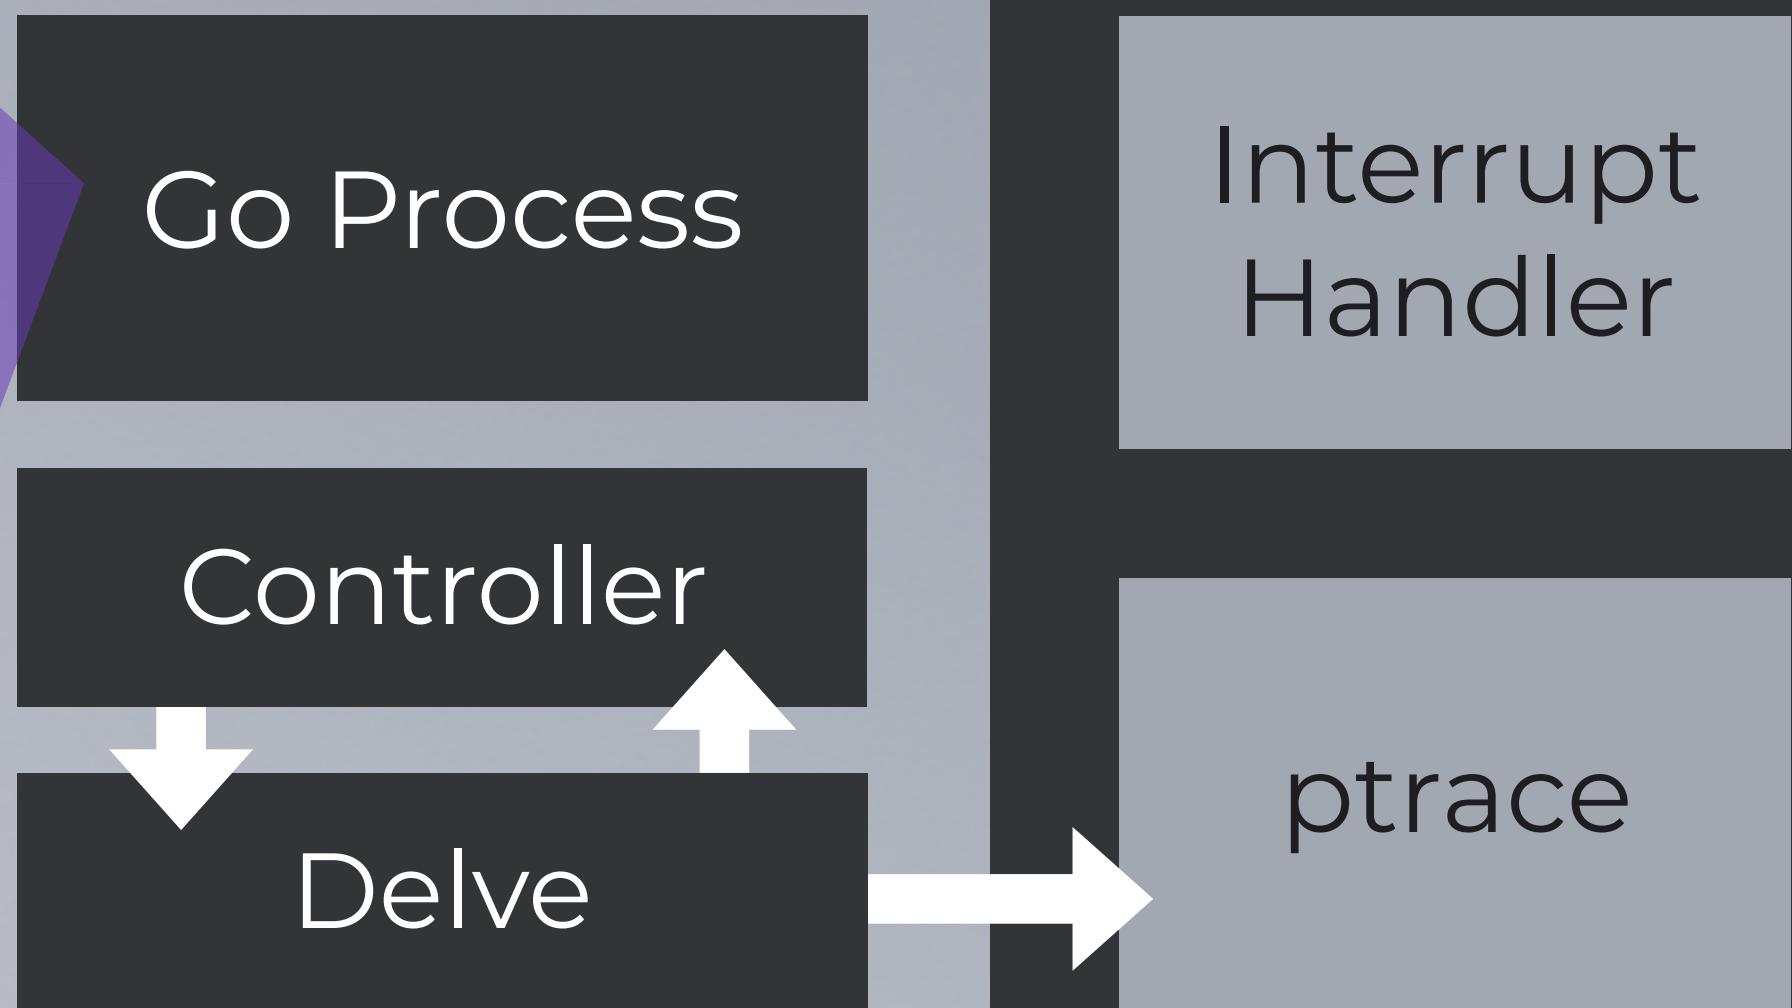
Controller

Delve

KERNEL

Interrupt
Handler

ptrace



00

B8

01

00

B8

00

B8

USER

Go Process

KERNEL

Interrupt
Handler

00

B8

01

00

B8

00

B8

USER

KERNEL

Go Process

Interrupt
Handler

00

B8

01

00

B8

00

B8

USER

KERNEL

Go Process

Interrupt
Handler

uprobes+BPF

uprobes

uprobes allows you to trace any instruction in user land with much less overhead than ptrace.

**Go Process
Running as Normal**

00

B8

01

00

B8

00

B8

USER

Go Process

KERNEL

uprobe
Handler

SSOL
Buffer



The Kernel Sets a Breakpoint

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

uprobe
Handler

B8

00

SSOL
Buffer

**When This is Hit the
uprobe Handler Fires**

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

uprobe
Handler

B8

00

SSOL
Buffer

00

B8

01

00

CC

00

B8

USER

KERNEL

Go Process

uprobe
Handler

B8

00

SSOL
Buffer



After the Handler Runs, The Copied Instruction is Single Stepped and the Program Resumes

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

uprobe
Handler

B8

00

SSOL
Buffer

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

uprobe
Handler

B8

00

SSOL
Buffer



00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

uprobe
Handler

B8

00

SSOL
Buffer



How do you write a
uprobe handler?

Write a kernel module?
A custom kernel?

BPF

BPF is a custom instruction set that you can use to build programs and inject them into the kernel.

The kernel validates the program to make sure it is safe and then compiles it for your architecture so it runs fast.

You can then attach these programs to various events, including uprobes.

BCC

BCC is a compiler for BPF programs written in C

It assists with interacting with your BPF programs

It is implemented as a library (libbcc.so)

**Go Process
Running as Normal**

00

B8

01

00

B8

00

B8

USER

KERNEL

Go Process



BCC Compiles and Injects BPF Program

00

B8

01

00

B8

00

B8

USER

KERNEL

Go Process

BCC

00

B8

01

00

B8

00

B8

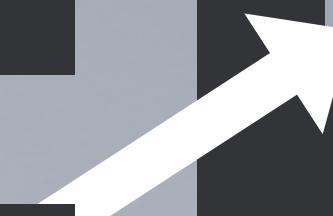
USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler



**BCC Creates uprobe and
Attaches it to the BPF Program**

00

B8

01

00

B8

00

B8

USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler

uprobe

00

B8

01

00

CC

00

B8

USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler

B8

00

uprobe

**When uprobe is Hit the
BPF Program is Ran**

00

B8

01

00

CC

00

B8

USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler

B8

00

uprobe

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

BPF
uprobe
Handler

B8
00

uprobe



BCC

00

B8

01

00

CC

00

B8

USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler

B8

00

uprobe

00

B8

01

00

CC

00

B8

USER

Go Process

KERNEL

BPF
uprobe
Handler

B8
00

uprobe



00

B8

01

00

CC

00

B8

USER

Go Process

BCC

KERNEL

BPF
uprobe
Handler

B8
00

uprobe

**Asynchronously Data is Shared
with the BCC Program**

00

B8

01

00

CC

00

B8

USER

KERNEL

Go Process

BCC

BPF
uprobe
Handler

B8

00

uprobe



bpftrace

Let's see an example of using bpftrace

```
uprobe:/tmp/counter:"main.doWork" {
    @ = *uaddr( "main.counter" );
}

interval:s:1 {
    printf( "counter: %d\t %d ops/sec\n", @, @-@prev);
    @prev = @;
}
```

```
$ ./tmp/counter --report  
This is a blackbox. Read my counter.  
counter: 32,350 (32,350 ops/s)
```

```
$ ./global.bt
```

It would be much better if we could write these programs in Go

A color photograph of a man with short brown hair, wearing a dark suit jacket, a light-colored shirt, and a patterned tie. He is positioned in front of a large, vertical film reel. The film reel has several frames visible, showing a scene from a movie. The background is dark and out of focus.

H
HD
HISTORY.COM

goBPF

goBPF provides Go bindings for libbcc

You have to write your BPF programs in C

Let's see an example of using goBPF


```
const bpfSource = `

// declare an array to share data between kernel and user space
BPF_ARRAY(count, u64, 1);

// define the function to handle the uprobe events
int read_counter()
{
    // get pointer to counter in program's memory
    u64 *counterPtr = (u64 *)%d;

    // get pointer to array in kernel memory
    int first = 0;
    u64 zero = 0, *val;
    val = count.lookup_or_init(&first, &zero);

    // copy value from program's memory to kernel memory
    bpf_probe_read(val, sizeof(*counterPtr), counterPtr);
}`
```

```
const bpfSource = `

// declare an array to share data between kernel and user space
BPF_ARRAY(count, u64, 1);

// define the function to handle the uprobe events
int read_counter()
{
    // get pointer to counter in program's memory
    u64 *counterPtr = (u64 *)%d;

    // get pointer to array in kernel memory
    int first = 0;
    u64 zero = 0, *val;
    val = count.lookup_or_init(&first, &zero);

    // copy value from program's memory to kernel memory
    bpf_probe_read(val, sizeof(*counterPtr), counterPtr);
}`
```

```
const bpfSource = `

// declare an array to share data between kernel and user space
BPF_ARRAY(count, u64, 1);

// define the function to handle the uprobe events
int read_counter()
{
    // get pointer to counter in program's memory
    u64 *counterPtr = (u64 *)%d;

    // get pointer to array in kernel memory
    int first = 0;
    u64 zero = 0, *val;
    val = count.lookup_or_init(&first, &zero);

    // copy value from program's memory to kernel memory
    bpf_probe_read(val, sizeof(*counterPtr), counterPtr);
}`
```

```
const bpfSource = `

// declare an array to share data between kernel and user space
BPF_ARRAY(count, u64, 1);

// define the function to handle the uprobe events
int read_counter()
{
    // get pointer to counter in program's memory
    u64 *counterPtr = (u64 *)%d;

    // get pointer to array in kernel memory
    int first = 0;
    u64 zero = 0, *val;
    val = count.lookup_or_init(&first, &zero);

    // copy value from program's memory to kernel memory
    bpf_probe_read(val, sizeof(*counterPtr), counterPtr);
}`
```

```
const bpfSource = `

// declare an array to share data between kernel and user space
BPF_ARRAY(count, u64, 1);

// define the function to handle the uprobe events
int read_counter()
{
    // get pointer to counter in program's memory
    u64 *counterPtr = (u64 *)%d;

    // get pointer to array in kernel memory
    int first = 0;
    u64 zero = 0, *val;
    val = count.lookup_or_init(&first, &zero);

    // copy value from program's memory to kernel memory
    bpf_probe_read(val, sizeof(*counterPtr), counterPtr);
}`
```

```
counterAddr := lookupSym( "/tmp/counter" , "main.counter" )
```

```
func lookupSym(path, name string) uint64 {
    f, _ := elf.Open(path)
    syms, _ := f.Symbols()
    for _, s := range syms {
        if s.Name == name {
            return s.Value
        }
    }
    log.Fatalf("Unable to find sym: %s %s", path, name)
    return 0
}
```

```
m := bpf.NewModule(fmt.Sprintf(bpfSource, counterAddr), nil)
probe, _ := m.LoadUprobe("read_counter")
m.AttachUprobe("/tmp/counter", "main.doWork", probe, -1)
table := bpf.NewTable(m.TableId("count"), m)
```

```
var prev, count uint64
for {
    data, _ := table.Get([]byte{0})
    count = binary.LittleEndian.Uint64(data)
    fmt.Printf("counter: %d\t(%d ops/s)\n", count, count-prev)
    prev = count
    time.Sleep(time.Second)
}
```

```
$ /tmp/counter --report
```

⋮

```
$ /tmp/gobpf
```

We have three really great options:

delve bpftrace goBPF

	delve	bpftrace	goBPF
intercept any instruction	✓	✓	✓
understands go	✓	✗	✗
r/w memory	✓	✓	✓
r/w args/return vals	✓	✗	✓
read registers	✓	✓	✓
compute latencies	✓	✗	✗
can use go packages	✓	✗	✓
high frequency events	✗	✓	✓

Static instrumentation is great but does not provide total coverage of your programs.

Dynamic instrumentation allows you to observe every aspect of your program, as it runs, without it knowing.

There is no free lunch.

Delve is awesome and has a great API.

bpftrace and goBPF are excellent tools who's support for Go is continually improving.

With these new kinds of instrumentation we can achieve new levels of understanding about our programs!

I hope you find them as useful as I do!

Thank you!

wat.io

@jasonkeene

wat.io

@jasonkeene

wat.io/src

github.com/go-delve/delve

github.com/iovisor/bpftrace

github.com/iovisor/gobpf

github.com/iovisor/bcc